

```
In [1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
```

```
In [3]: from sklearn.datasets import fetch_openml
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_squared_error, mean_absolute_error, r2_s
```

```
In [5]: sns.set(style="whitegrid")
```

```
In [7]: boston = fetch_openml(name='boston', version=1, as_frame=True)
df = boston.frame

df.head()
```

```
Out[7]:
```

	CRIM	ZN	INDUS	CHAS	NOX	RM	AGE	DIS	RAD	TAX	PTRATIO	B	LSTAT
0	0.00632	18.0	2.31	0	0.538	6.575	65.2	4.0900	1	296.0	15.3	396.9	4.9875
1	0.02731	0.0	7.07	0	0.469	6.421	78.9	4.9671	2	242.0	17.8	397.0	4.9783
2	0.02729	0.0	7.07	0	0.469	7.185	61.1	4.9671	2	242.0	17.8	397.0	4.9783
3	0.03237	0.0	2.18	0	0.458	6.998	45.8	6.0622	3	222.0	18.7	397.0	4.9783
4	0.06905	0.0	2.18	0	0.458	7.147	54.2	6.0622	3	222.0	18.7	397.0	4.9783

◀ ▶

```
In [9]: df.shape
df.info()
df.describe()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 506 entries, 0 to 505
Data columns (total 14 columns):
 #   Column   Non-Null Count  Dtype  
 --- 
 0   CRIM      506 non-null    float64
 1   ZN        506 non-null    float64
 2   INDUS     506 non-null    float64
 3   CHAS      506 non-null    category
 4   NOX       506 non-null    float64
 5   RM         506 non-null    float64
 6   AGE        506 non-null    float64
 7   DIS         506 non-null    float64
 8   RAD         506 non-null    category
 9   TAX        506 non-null    float64
 10  PTRATIO    506 non-null    float64
 11  B          506 non-null    float64
 12  LSTAT      506 non-null    float64
 13  MEDV       506 non-null    float64
dtypes: category(2), float64(12)
memory usage: 49.0 KB
```

Out[9]:

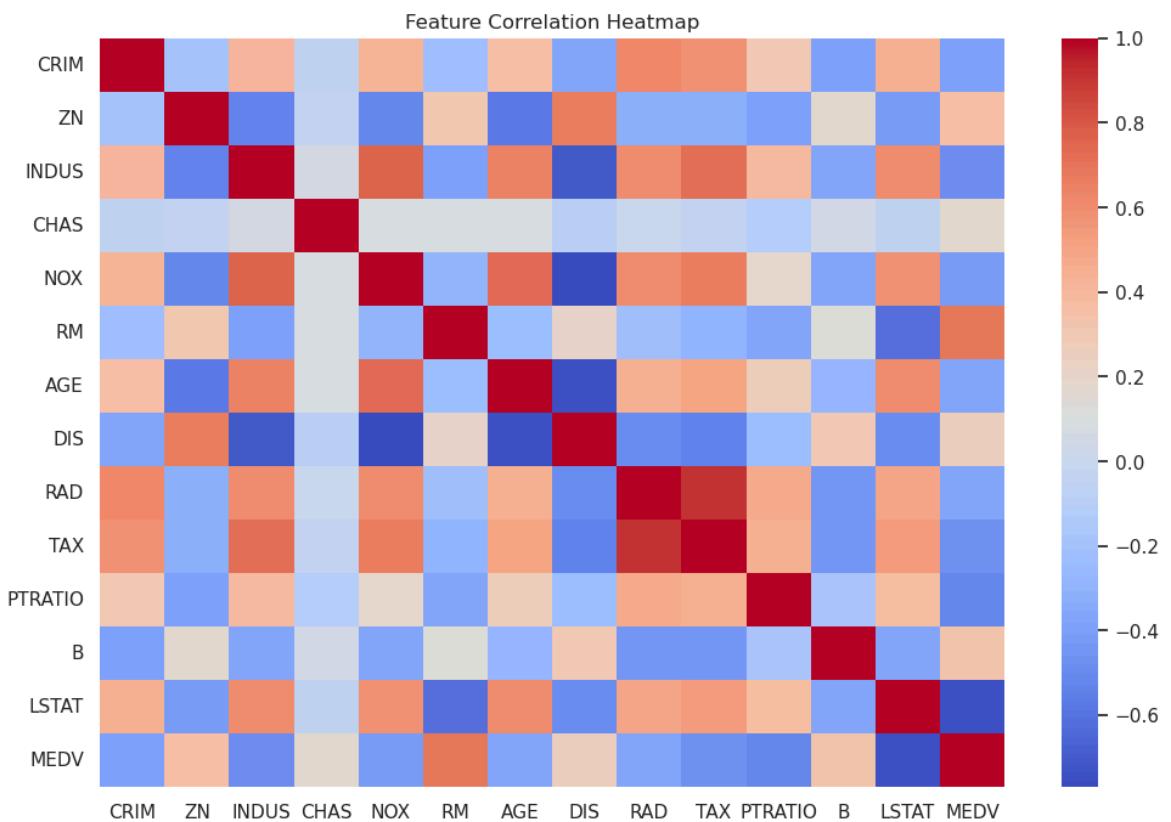
	CRIM	ZN	INDUS	NOX	RM	AGE	
<b>count</b>	506.000000	506.000000	506.000000	506.000000	506.000000	506.000000	506.000000
<b>mean</b>	3.613524	11.363636	11.136779	0.554695	6.284634	68.574901	3.7
<b>std</b>	8.601545	23.322453	6.860353	0.115878	0.702617	28.148861	2.1
<b>min</b>	0.006320	0.000000	0.460000	0.385000	3.561000	2.900000	1.1
<b>25%</b>	0.082045	0.000000	5.190000	0.449000	5.885500	45.025000	2.1
<b>50%</b>	0.256510	0.000000	9.690000	0.538000	6.208500	77.500000	3.2
<b>75%</b>	3.677083	12.500000	18.100000	0.624000	6.623500	94.075000	5.1
<b>max</b>	88.976200	100.000000	27.740000	0.871000	8.780000	100.000000	12.1

In [11]: `df.isnull().sum()`

```
Out[11]: CRIM      0
          ZN       0
          INDUS    0
          CHAS     0
          NOX      0
          RM       0
          AGE      0
          DIS      0
          RAD      0
          TAX      0
          PTRATIO   0
          B        0
          LSTAT    0
          MEDV     0
          dtype: int64
```

In [13]: 

```
plt.figure(figsize=(12,8))
sns.heatmap(df.corr(), annot=False, cmap='coolwarm')
plt.title("Feature Correlation Heatmap")
plt.show()
```



```
In [15]: plt.figure(figsize=(6,4))
sns.scatterplot(x=df['RM'], y=df['MEDV'])
plt.title("Rooms vs Price")
plt.show()
```



```
In [17]: X = df.drop('MEDV', axis=1) # independent variables
y = df['MEDV'] # target variable
```

```
In [19]: X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.2, random_state=42)
```

```
)
```

```
print("Training samples:", X_train.shape)
print("Testing samples:", X_test.shape)
```

```
Training samples: (404, 13)
Testing samples: (102, 13)
```

```
In [21]: model = LinearRegression()
model.fit(X_train, y_train)
```

```
Out[21]: ▾ LinearRegression ⓘ ?
```

```
LinearRegression()
```

```
In [23]: coeff_df = pd.DataFrame(model.coef_, X.columns, columns=['Coefficient'])
coeff_df
```

```
Out[23]:
```

	Coefficient
CRIM	-0.113056
ZN	0.030110
INDUS	0.040381
CHAS	2.784438
NOX	-17.202633
RM	4.438835
AGE	-0.006296
DIS	-1.447865
RAD	0.262430
TAX	-0.010647
PTRATIO	-0.915456
B	0.012351
LSTAT	-0.508571

```
In [27]: X.dtypes
```

```
Out[27]: CRIM      float64
ZN        float64
INDUS     float64
CHAS      category
NOX       float64
RM        float64
AGE       float64
DIS        float64
RAD        category
TAX       float64
PTRATIO    float64
B         float64
LSTAT     float64
dtype: object
```

```
In [29]: X = X.apply(pd.to_numeric)
y = pd.to_numeric(y)
```

```
In [31]: from sklearn.model_selection import train_test_split

X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.2, random_state=42
)
```

```
In [33]: from sklearn.linear_model import LinearRegression

model = LinearRegression()
model.fit(X_train, y_train)
```

```
Out[33]: ▾ LinearRegression ⓘ ?
```

```
LinearRegression()
```

```
In [35]: y_pred = model.predict(X_test)
```

```
In [37]: from sklearn.metrics import mean_squared_error, r2_score
import numpy as np

print("RMSE:", np.sqrt(mean_squared_error(y_test, y_pred)))
print("R2 Score:", r2_score(y_test, y_pred))
```

RMSE: 4.928602182665346

R2 Score: 0.6687594935356307

```
In [39]: mae = mean_absolute_error(y_test, y_pred)
mse = mean_squared_error(y_test, y_pred)
rmse = np.sqrt(mse)
r2 = r2_score(y_test, y_pred)

print("Mean Absolute Error:", mae)
print("Mean Squared Error:", mse)
print("Root Mean Squared Error:", rmse)
print("R2 Score:", r2)
```


Mean Absolute Error: 3.189091965887852

Mean Squared Error: 24.291119474973613


Root Mean Squared Error: 4.928602182665346

R<sup>2</sup> Score: 0.6687594935356307

```
In [41]: plt.figure(figsize=(6,4))
plt.scatter(y_test, y_pred)
plt.xlabel("Actual Prices")
plt.ylabel("Predicted Prices")
plt.title("Actual vs Predicted Prices")
plt.show()
```



```
In [43]: residuals = y_test - y_pred  
  
plt.figure(figsize=(6,4))  
sns.histplot(residuals, kde=True)  
plt.title("Residual Distribution")  
plt.show()
```



```
In [ ]:
```