

```
In [1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
```

```
In [3]: from sklearn.datasets import fetch_openml
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_squared_error, mean_absolute_error, r2_s
```

```
In [5]: sns.set(style="whitegrid")
```

```
In [122...]: boston = fetch_openml(name='boston', version=1, as_frame=True)
#df1= pd.read_csv("Desktop/datasets/boston/train.csv")

df = boston.frame

df.head()
```

Out[122...]

	CRIM	ZN	INDUS	CHAS	NOX	RM	AGE	DIS	RAD	TAX	PTRATIO	B	LSTAT
0	0.00632	18.0	2.31	0	0.538	6.575	65.2	4.0900	1	296.0	15.3	396.9	4.98
1	0.02731	0.0	7.07	0	0.469	6.421	78.9	4.9671	2	242.0	17.8	397.0	4.98
2	0.02729	0.0	7.07	0	0.469	7.185	61.1	4.9671	2	242.0	17.8	397.0	4.98
3	0.03237	0.0	2.18	0	0.458	6.998	45.8	6.0622	3	222.0	18.7	397.0	4.98
4	0.06905	0.0	2.18	0	0.458	7.147	54.2	6.0622	3	222.0	18.7	397.0	4.98

◀ ▶

```
In [86]: df.shape
df.info()
df.describe()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 506 entries, 0 to 505
Data columns (total 14 columns):
 #   Column      Non-Null Count  Dtype  
 --- 
 0   CRIM        506 non-null    float64
 1   ZN          506 non-null    float64
 2   INDUS       506 non-null    float64
 3   CHAS         506 non-null    category
 4   NOX          506 non-null    float64
 5   RM           506 non-null    float64
 6   AGE          506 non-null    float64
 7   DIS           506 non-null    float64
 8   RAD           506 non-null    category
 9   TAX           506 non-null    float64
 10  PTRATIO      506 non-null    float64
 11  B             506 non-null    float64
 12  LSTAT        506 non-null    float64
 13  MEDV          506 non-null    float64
dtypes: category(2), float64(12)
memory usage: 49.0 KB
```

Out[86]:

	CRIM	ZN	INDUS	NOX	RM	AGE	
count	506.000000	506.000000	506.000000	506.000000	506.000000	506.000000	506.000000
mean	3.613524	11.363636	11.136779	0.554695	6.284634	68.574901	3.7
std	8.601545	23.322453	6.860353	0.115878	0.702617	28.148861	2.1
min	0.006320	0.000000	0.460000	0.385000	3.561000	2.900000	1.1
25%	0.082045	0.000000	5.190000	0.449000	5.885500	45.025000	2.1
50%	0.256510	0.000000	9.690000	0.538000	6.208500	77.500000	3.2
75%	3.677083	12.500000	18.100000	0.624000	6.623500	94.075000	5.1
max	88.976200	100.000000	27.740000	0.871000	8.780000	100.000000	12.1



In [88]:

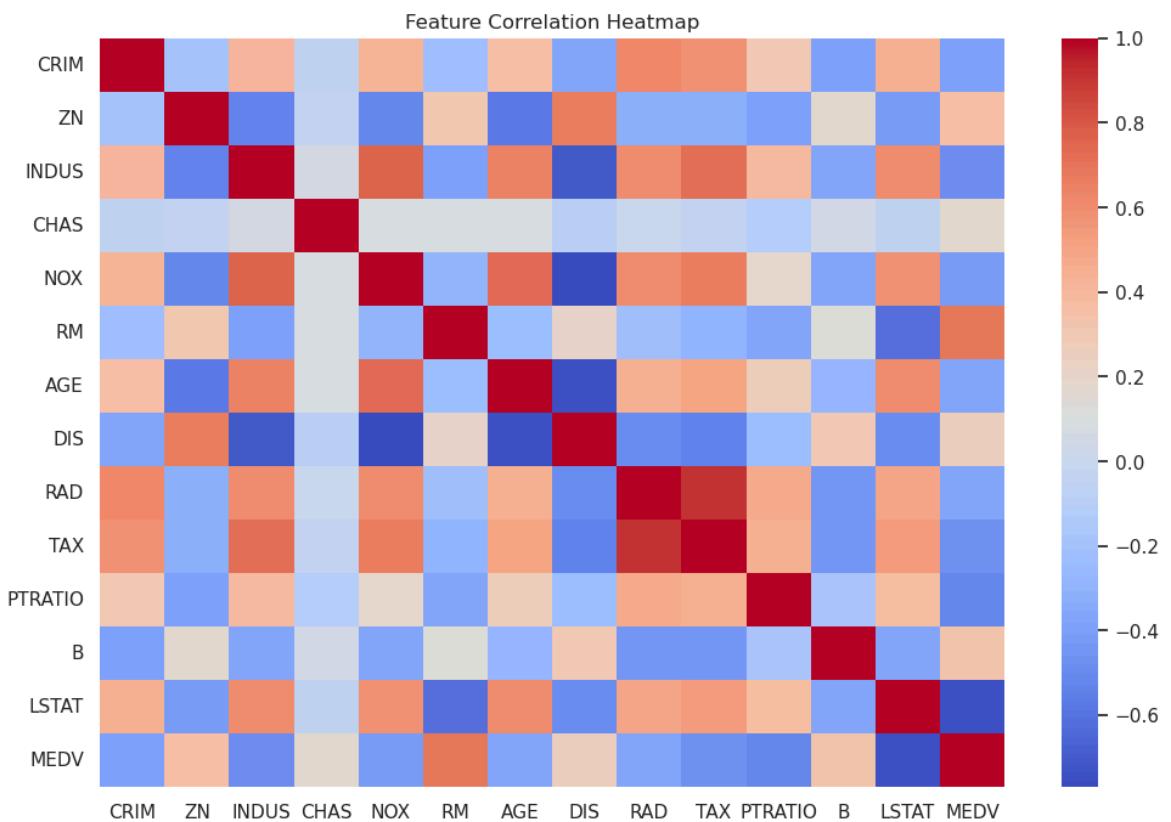
```
df.isnull().sum()
```

Out[88]:

```
CRIM      0
ZN        0
INDUS     0
CHAS      0
NOX       0
RM        0
AGE       0
DIS       0
RAD       0
TAX       0
PTRATIO   0
B          0
LSTAT     0
MEDV      0
dtype: int64
```

In [90]:

```
plt.figure(figsize=(12,8))
sns.heatmap(df.corr(), annot=False, cmap='coolwarm')
plt.title("Feature Correlation Heatmap")
plt.show()
```



```
In [92]: plt.figure(figsize=(6,4))
sns.scatterplot(x=df['RM'], y=df['MEDV'])
plt.title("Rooms vs Price")
plt.show()
```



```
In [94]: X = df.drop('MEDV', axis=1) # independent variables
y = df['MEDV'] # target variable
```

```
In [96]: X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.2, random_state=42)
```

```
)
```

```
print("Training samples:", X_train.shape)
print("Testing samples:", X_test.shape)
```

```
Training samples: (404, 13)
Testing samples: (102, 13)
```

```
In [98]: model = LinearRegression()
model.fit(X_train, y_train)
```

```
Out[98]:
```

▼ LinearRegression ⓘ ⓘ

LinearRegression()

```
In [100... coeff_df = pd.DataFrame(model.coef_, X.columns, columns=['Coefficient'])
coeff_df
```

```
Out[100...          Coefficient
```

	Coefficient
CRIM	-0.113056
ZN	0.030110
INDUS	0.040381
CHAS	2.784438
NOX	-17.202633
RM	4.438835
AGE	-0.006296
DIS	-1.447865
RAD	0.262430
TAX	-0.010647
PTRATIO	-0.915456
B	0.012351
LSTAT	-0.508571

```
In [102... X.dtypes
```

```
Out[102... CRIM      float64
ZN        float64
INDUS     float64
CHAS      category
NOX       float64
RM        float64
AGE       float64
DIS        float64
RAD        category
TAX       float64
PTRATIO    float64
B         float64
LSTAT     float64
dtype: object
```

```
In [104... X = X.apply(pd.to_numeric)
      y = pd.to_numeric(y)
```

```
In [106... from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.2, random_state=42
)
```

```
In [108... from sklearn.linear_model import LinearRegression
model = LinearRegression()
model.fit(X_train, y_train)
```

```
Out[108... ▾ LinearRegression ⓘ ?]
LinearRegression()
```

```
In [110... y_pred = model.predict(X_test)
```

```
In [112... from sklearn.metrics import mean_squared_error, r2_score
import numpy as np
print("RMSE:", np.sqrt(mean_squared_error(y_test, y_pred)))
print("R2 Score:", r2_score(y_test, y_pred))
```

RMSE: 4.928602182665346

R2 Score: 0.6687594935356307

```
In [114... mae = mean_absolute_error(y_test, y_pred)
mse = mean_squared_error(y_test, y_pred)
rmse = np.sqrt(mse)
r2 = r2_score(y_test, y_pred)

print("Mean Absolute Error:", mae)
print("Mean Squared Error:", mse)
print("Root Mean Squared Error:", rmse)
print("R² Score:", r2)
```

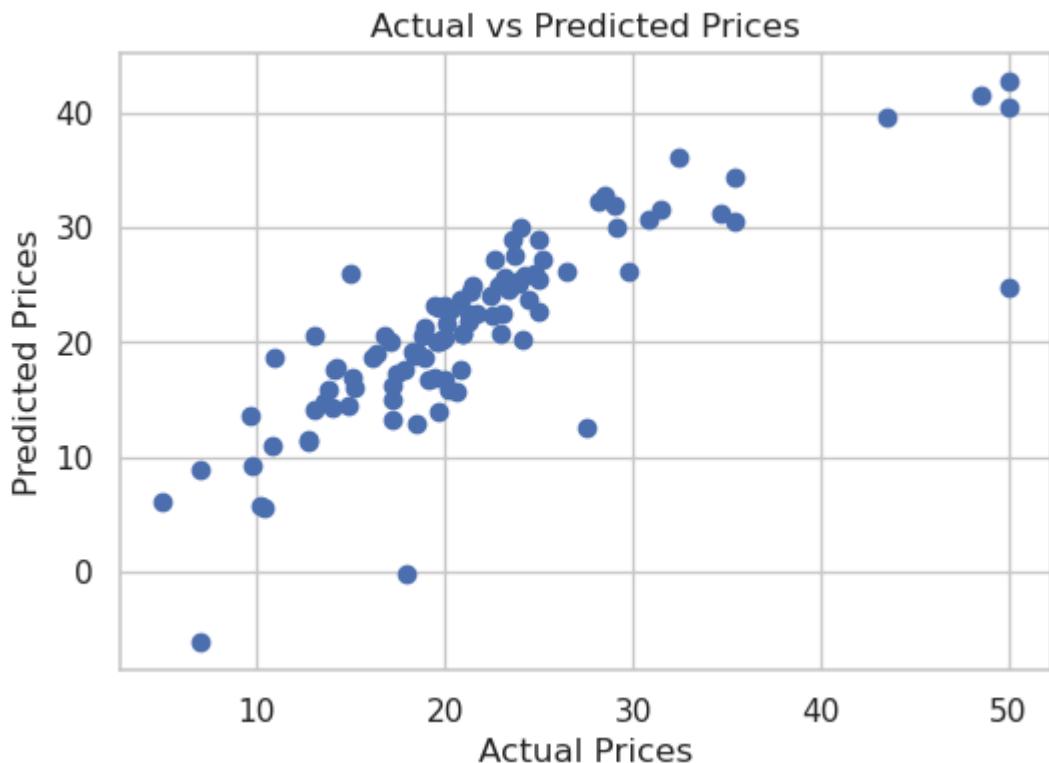
Mean Absolute Error: 3.189091965887852

Mean Squared Error: 24.291119474973613

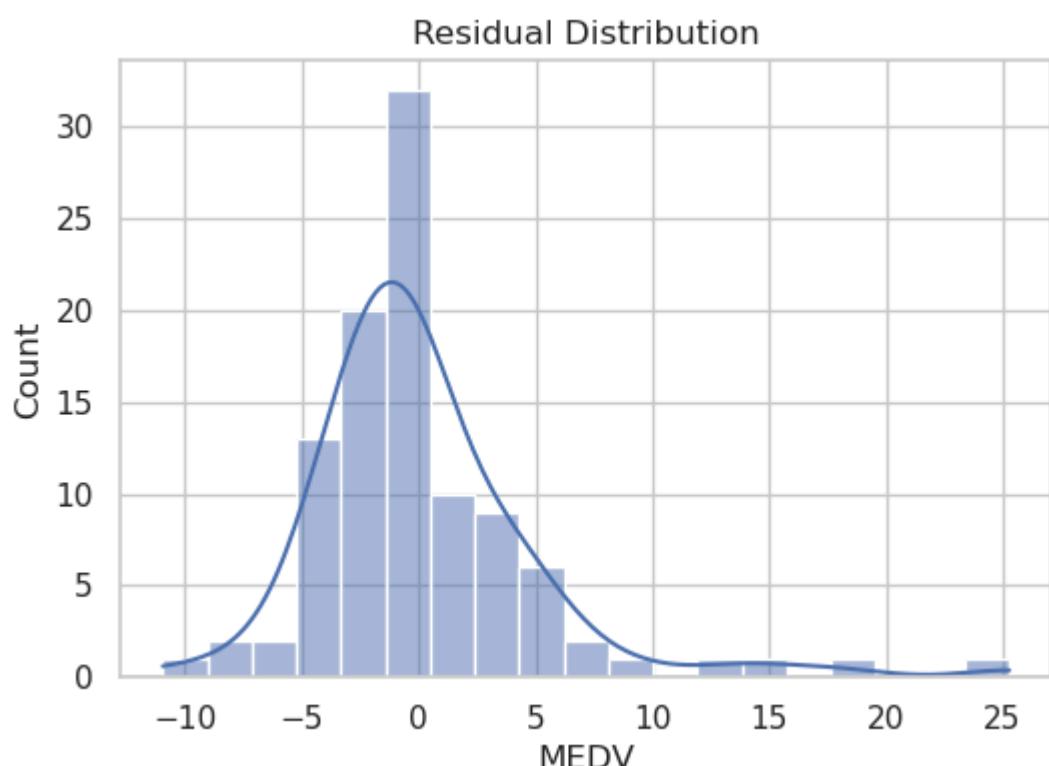
Root Mean Squared Error: 4.928602182665346

R² Score: 0.6687594935356307

```
In [116... plt.figure(figsize=(6,4))
plt.scatter(y_test, y_pred)
plt.xlabel("Actual Prices")
plt.ylabel("Predicted Prices")
plt.title("Actual vs Predicted Prices")
plt.show()
```



```
In [118]: residuals = y_test - y_pred  
  
plt.figure(figsize=(6,4))  
sns.histplot(residuals, kde=True)  
plt.title("Residual Distribution")  
plt.show()
```



```
In [ ]: ##multiple
```

```
In [124]: import pandas as pd  
import numpy as np
```

```
import matplotlib.pyplot as plt
import seaborn as sns

from sklearn.datasets import fetch_openml
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_absolute_error, mean_squared_error, r2_s
```

In [126...]
boston = fetch_openml(name='boston', version=1, as_frame=True)
df = boston.frame
df = df.apply(pd.to_numeric) # ensure all values are numeric

In [128...]
print(df.shape)
df.head()

(506, 14)

Out[128...]

	CRIM	ZN	INDUS	CHAS	NOX	RM	AGE	DIS	RAD	TAX	PTRATIO	B
0	0.00632	18.0	2.31	0	0.538	6.575	65.2	4.0900	1	296.0	15.3	39.6
1	0.02731	0.0	7.07	0	0.469	6.421	78.9	4.9671	2	242.0	17.8	39.6
2	0.02729	0.0	7.07	0	0.469	7.185	61.1	4.9671	2	242.0	17.8	39.6
3	0.03237	0.0	2.18	0	0.458	6.998	45.8	6.0622	3	222.0	18.7	39.6
4	0.06905	0.0	2.18	0	0.458	7.147	54.2	6.0622	3	222.0	18.7	39.6

◀ ▶

In [130...]
X = df.drop("MEDV", axis=1) # multiple input variables
y = df["MEDV"] # house price

In [132...]
X_train, X_test, y_train, y_test = train_test_split(
 X, y, test_size=0.2, random_state=42
)

In [134...]
model = LinearRegression()
model.fit(X_train, y_train)

Out[134...]
▼ LinearRegression ⓘ ?
LinearRegression()

In [136...]
coeff_df = pd.DataFrame(model.coef_, X.columns, columns=['Coefficient'])
coeff_df

Out[136...]

	Coefficient
CRIM	-0.113056
ZN	0.030110
INDUS	0.040381
CHAS	2.784438
NOX	-17.202633
RM	4.438835
AGE	-0.006296
DIS	-1.447865
RAD	0.262430
TAX	-0.010647
PTRATIO	-0.915456
B	0.012351
LSTAT	-0.508571

In [138...]

```
y_pred = model.predict(X_test)
```

In [140...]

```
mae = mean_absolute_error(y_test, y_pred)
mse = mean_squared_error(y_test, y_pred)
rmse = np.sqrt(mse)
r2 = r2_score(y_test, y_pred)

print("MAE:", mae)
print("MSE:", mse)
print("RMSE:", rmse)
print("R2 Score:", r2)
```

MAE: 3.189091965887852
MSE: 24.291119474973613
RMSE: 4.928602182665346
R² Score: 0.6687594935356307

In [142...]

```
plt.scatter(y_test, y_pred)
plt.xlabel("Actual Prices")
plt.ylabel("Predicted Prices")
plt.title("Actual vs Predicted Prices")
plt.show()
```



```
In [144...]: #  
  
In [146...]: #singular  
  
In [148...]: import pandas as pd  
import numpy as np  
import matplotlib.pyplot as plt  
  
from sklearn.datasets import fetch_openml  
from sklearn.model_selection import train_test_split  
from sklearn.linear_model import LinearRegression  
from sklearn.metrics import mean_squared_error, r2_score  
  
In [150...]: boston = fetch_openml(name='boston', version=1, as_frame=True)  
df = boston.frame  
df = df.apply(pd.to_numeric)  
  
In [152...]: X = df[["RM"]]      # singular feature  
y = df["MEDV"]        # target  
  
In [154...]: X_train, X_test, y_train, y_test = train_test_split(  
           X, y, test_size=0.2, random_state=42  
         )  
  
In [156...]: model = LinearRegression()  
model.fit(X_train, y_train)
```

Out[156...]

```
LinearRegression
```

```
LinearRegression()
```

In [162...]

```
print("Slope (Coefficient):", model.coef_[0])
print("Intercept:", model.intercept_)
#Price=(Slope×RM)+Intercept
```

Slope (Coefficient): 9.348301406497727
Intercept: -36.24631889813795

In [164...]

```
y_pred = model.predict(X_test)
```

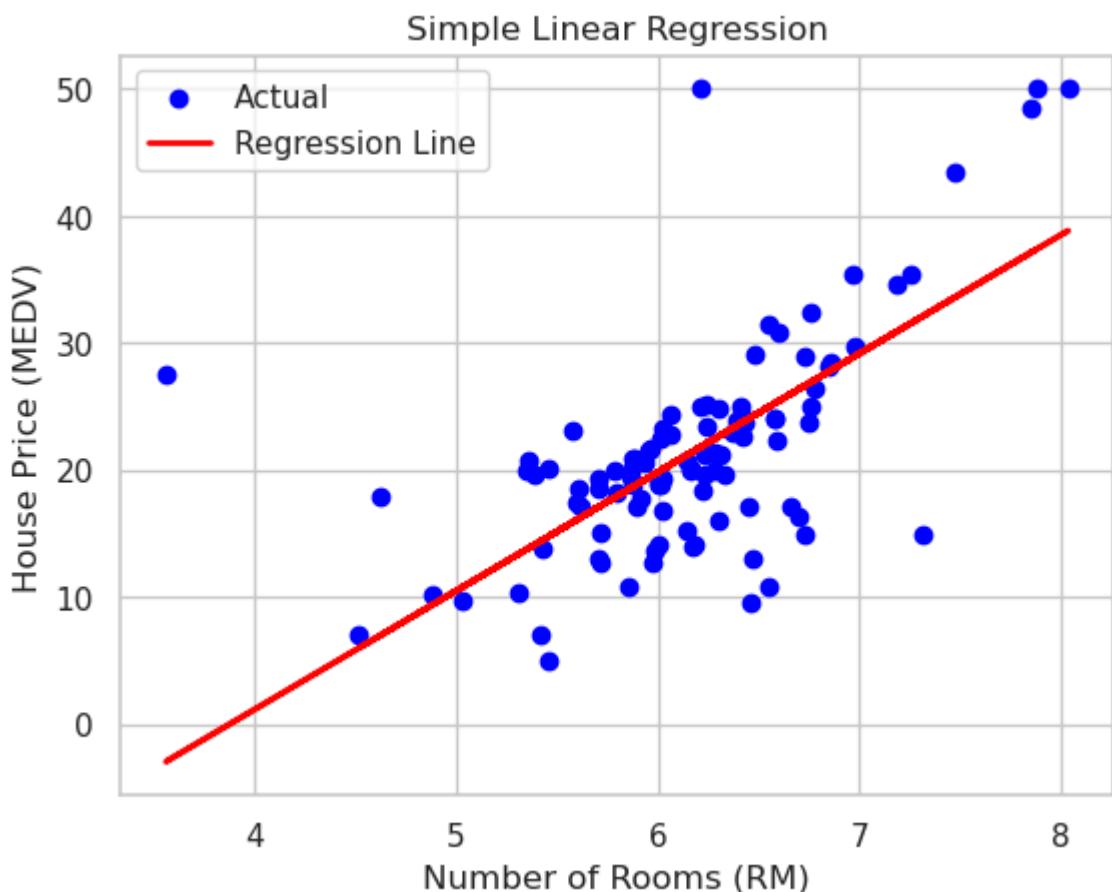
In [166...]

```
print("RMSE:", np.sqrt(mean_squared_error(y_test, y_pred)))
print("R2 Score:", r2_score(y_test, y_pred))
```

RMSE: 6.792994578778734
R2 Score: 0.3707569232254778

In [168...]

```
plt.scatter(X_test, y_test, color='blue', label="Actual")
plt.plot(X_test, y_pred, color='red', linewidth=2, label="Regression Line")
plt.xlabel("Number of Rooms (RM)")
plt.ylabel("House Price (MEDV)")
plt.title("Simple Linear Regression")
plt.legend()
plt.show()
```



In []: